Convolutional Neural Networks for Classifying Handwritten Digits

Favyen Bastani
fbastani@perennate.com

Zhengli Wang
wzl@mit.edu

December 5, 2014

Abstract

We implement a CPU-based convolutional neural network (CNN) library in Python, using the `numpy` library for fast array operations. Our library constructs the complex structure of the CNN, performs forward propagation and pooling, and trains the network using backward propagation and stochastic gradient descent. We evaluate CNN performance in terms of classification accuracy and training time on the MNIST database under a wide range of network parameters, and compare CNN to feature extraction approaches such as HOG. We find that CNN is able to achieve much lower misclassification rate (as low as 1.73%), but requires significantly more time for training of features.

1 Introduction

Image recognition algorithms have become ubiquitous: on Google Maps Street View, image recognition protects privacy by automatically blurring faces and license plate numbers [2]; robots use object recognition to autonomously move or assemble structures; meanwhile, individuals routinely employ optical character recognition for scanning documents and facial recognition for labeling photographs. Over the last decade, image recognition approaches have developed dramatically. Initially, feature extraction techniques such as HOG and FAST were typically combined with support vector machine (SVM) classifiers, with research often focusing on selection of SVM kernels and defining features for specific image types. Recently, though, neural networks have received renewed attention and achieved high classification performance.

Cells in the mammalian visual cortex are able to identify local patterns in the visual field [5]. Convolutional neural networks (CNNs) aim to do the same through *convolutional and pooling layers*. While traditional neural networks are composed of densely connected layers, where each neu-

ron in the previous layer connects to each neuron in the next layer, convolutional layers restrict filters to be applied on local sub-regions of an input image. By training these layers through backpropagation, these filters eventually detect specific features in the image that help with classification. CNNs have achieved high image recognition accuracy on a large number of datasets, including ImageNet (for object recognition) [4], face recognition [7], and video quality assessment [8].

In this project, we implement a fast CPU-based CNN toolkit [1] in Python, and apply it to the MNIST handwritten digits database [9]. We evaluate the performance of various neural network parameters (including layer configuration, output layer activation function, and training algorithm) in terms of classification accuracy and training speed, and compare CNN to feature extraction approaches. We find that while CNNs require significantly more training time, they outperform other algorithms with minimal tuning of the network.

In Sections 2 and 3, we detail convolution and pooling functions, the convolutional neural network structure and training algorithm, and our implementation. In Section 5, we experiment with various neural network parameters, and in Section 6, we compare with SVM-based algorithms. Finally, we conclude in Section 7.

2 Neural networks

An artificial neural network (ANN) consists of neuron units that, when composed in a layered network, compute some output function from an input vector $y = h(x)$. Each neuron in the network has a specific set of input values (which may be elements of the input vector of the network, or outputs from other neurons) and a single output. The output is generally computed as $n_{w,b}(z) = f(w \cdot z + b)$ for some *activation function* f, weight vector w, and bias parameter b. The sigmoid function is often used as the activation function, i.e. $f(l) = \frac{1}{1+e^{-l}}$,

1

because it mimics the hard threshold activation pattern of a biological neuron.

Neurons are connected to form a neural network. Typically, the network is composed of layers of neurons, where neurons in the first layer accept inputs from the input vector x, and neurons in each following layer take inputs from the outputs of the neurons in the preceding layer. Then, the network output vector y consists of the outputs from neurons in the last layer of the network. In fully connected neural networks, a neuron takes inputs from every neuron in the previous layer.

Forward propagation refers to the process of propagating some input vector x through each layer of the neural network to compute $y = h(x)$. In a fully connected neural network, this can be done using efficient matrix operations (which can take advantage of fast matrix libraries such as `numpy`). Let Z_l be the output vector of the lth neural network layer (where $Z_0 = x$) and n_{li} be the ith neuron in layer l; then, $Z_{l+1} = f(W^T Z_l + B)$ where W_{ij} is the weight from n_{li} to $n_{l+1,j}$ and B_j is the bias parameter for neuron $n_{l+1,j}$.

Neural networks can be trained using back propagation. At a high level, we compute a cost function on the outputs from the neural network given some input and desired output, and propagate the error (in the form of partial derivatives of the cost) starting at the last layer to the first layer. These error terms are then used to compute the derivative with respect to specific weight and bias parameters, which can then be updated via gradient descent.

2.1 ANN for multi-label classification

Artificial neural networks are often used for multi-label classification. However, the sigmoid activation function is not well suited for this purpose, where we want to assign exactly one label c for each input vector x. Instead, we use *softmax activation function*, defined as

$$f(L, i) = \frac{e^{L_i}}{\sum_j e^{L_j}}$$

Here, L is a vector containing inputs to the activation function for each neuron in the softmax layer. For back propagation, the cost function given a desired output vector o is the cross-entropy error function [10],

$$E = -\sum_i o_i \log(y_i)$$

2.2 Stochastic gradient descent

In normal gradient descent neural network training, we iteratively update parameters (weights and bias terms) after averaging derivatives from back propagation across all training samples. However, when the training set is large, each iteration of gradient descent may take a significant amount of time to complete. Stochastic gradient descent instead randomly selects mini-batches of a predetermined size from the training samples on each iteration and updates parameters by averaging derivatives across only the selected samples (the batch size is generally selected to yield optimal performance with matrix libraries, although in our implementation we perform back propagation separately for each sample). This increases solution sparsity as well as convergence speed.

We implement one variation of stochastic gradient descent that further reduces training time in many cases. Rather than updating parameters directly from the computed gradient ∇, we apply a momentum that retains gradients from previous iterations. Let $v_0 = 0$ be the initial velocity. For a training rate α and momentum m, we update parameters θ_i on the ith iteration as [6]

$$v_i = mv_{i-1} + \alpha \nabla_i$$
$$\theta_i = \theta_{i-1} - v_i$$

Note that the training rate α is generally smaller in stochastic gradient descent than in conventional gradient descent because the random mini-batch selection introduces greater variance on the training process.

3 Convolutional neural networks

Images, along with many other data types such as audio signals, exhibit strong local patterns that need to be identified for classification. In some cases, fully connected neural networks may ignore these patterns, perhaps overfitting instead; in other cases, training fully connected networks to identify local patterns and achieve high classification accuracy takes a significant number of training iterations. Convolutional neural networks take an alternate approach, where the localized pattern properties are directly encoded into the structure of the neural network. In this section, we discuss the functionality of convolution and pooling layers, and strategies for incorporating them into neural networks.

3.1 Convolution

Images can often be hierarchically broken down into features and sub-features that support image recognition. For example, a flower may contain a stem and a petal, each of which consist of edges at various positions and angles. Additionally, these features may exist at various locations in the image: if an image is shifted in one direction, an image recognition algorithm should still be able to classify it in the same way.

The idea behind convolutional layers for image recognition is that a feature can be represented as an n by n *filter*, and then convolved two-dimensionally with the m by m input (which in this case would be an image). If regions of the input match up with the filter, then elements of the convolution output corresponding to those regions will generally take higher values. The output from the convolution is added with a bias term and passed through an activation function. Notably, each component of the convolution can be modeled as a neuron: each neuron takes inputs from a sub-region, multiplies the input by a weight matrix (the filter), and then applies the activation function; besides the sparsely connected input, the only other difference from conventional neural networks is that we require the weight matrix to be the same across multiple neurons (since convolution applies a single filter repeatedly across the image). This enables the detection of local patterns that we desire, since the filter is restricted to only consider sub-regions of the input, while still incorporating the weights and bias components that define neural networks.

In each convolutional layer, we have k filters, each an n by n matrix. Then, for an m by m input, we get a k by $(m - n + 1)$ by $(m - n + 1)$ output from the two-dimensional convolution.

3.2 Pooling

Pooling layers apply an aggregate function (typically either max or mean) on non-overlapping regions of the output from convolutional layers (this is done separately for each filter). These regions are often equal in size; so if the output from the convolutional layer is 30 by 30, the pooling layer may pool across each of the four 15 by 15 regions, yielding four output values.

Pooling accomplishes two purposes. First, we reduce the amount of data that later layers need to consider, speeding up training and reducing overfitting. Second, we get a degree of translation invariance. If the input is shifted slightly, the result after pooling will still be similar. This is desirable since generally image recognition classes should remain the same under basic transforms.

3.3 Network

Convolution and pooling can be used in isolation, without incorporation in a convolutional neural network. In this case, an autoencoder is typically used to train the filters, where the training goal is to have the output layer match the input layer despite a smaller number of neurons in intermediate layers. The trained filters can then be used to extract features from both training and testing data, and then fed into any classifier. This also supports unsupervised learning of features.

Convolutional neural networks present a more interesting case, where the convolution and pooling operations are combined with classification into a single structure that can be trained in one process. Typical convolutional neural networks consist of one to three convolutional layers followed by fully connected hidden layers; back propagation is carried out across the entire network. When there are several convolutional layers, the first convolutional layer often is used by the network as an edge detector (this determination happens from training alone, without any manual specification other than the layer sizes), while later convolutional layers in the network detect more and more complex features; for simple datasets, a single convolutional layer can be applied to directly identify features. The hidden layers use final pooling outputs to eventually classify the image.

4 Experimental setup

We implement a CPU-based convolutional neural network library [1] in Python using the `numpy` package, which provides fast scientific computing operations. Our implementation includes a `Network` class which has attributes that represent the number of layers and the number of neurons in each layer. Our code can use either "softmax" or "sigmoid" activation functions of neurons in the fully connected layers of the network. Our code initializes weights to be random and uniformly distributed within [-1,1], which serves the purpose of breaking symmetry so that filters train in different directions. The network class has separate functions to support forward propagation and backward propagation. For forward propagation, it simply involves looping over all the neurons, computing the output of the current layer from the output of the previous layer's values; the output from

Figure 1: Training samples from the MNIST dataset that show the difficulty of achieving high classification accuracies.

Parameter	Value
Filter sidelength n	13
Number of filters k	64
Pooling blocks	2 by 2
Pooling method	Mean
Output layer activation	Softmax
Training algorithm	SGD
Training duration	10 hours

Table 1: Default parameter values used for performance evaluation.

Each MNIST image contains a single handwritten digit (from zero to nine). The images are all 28 pixels by 28 pixels and grayscale. Figure 1 shows four training samples, two 4 digits and two 8 digits; since digits can be written in multiple ways and at different angles and sizes, classification is a difficult problem.

For our experimentation, we use a basic default setting of the neural network, and then adjust network parameters to identify patterns in the change in performance, such as test error rate and average per-iteration runtime. Specifically, we use a default setting of 64 filters, each 13 by 13, with a pooling layer of size 2 by 2 (meaning we have four pooling blocks, each pooling over an 8 by 8 region of the convolution output). The output from the pooling stage, then, consists of 256 values, which are fully connected to a final output layer consisting of ten neurons (without additional hidden layers in between); each output neuron corresponds to one digit, and the maximum value over the neurons is the classification output. Due to the long training times needed for CNN, we only show results from one trial (although we find results are relatively stable even when using stochastic gradient descent). Experiments are conducted on twenty-four thread machines with dual Intel Xeon L5640 2.27 GHz CPUs.

the final layer is returned. In backward propagation, our code computes and averages the derivatives across a subset of the data (the subset is chosen randomly per stochastic gradient descent algorithm); this step is the most computationally expensive. After that, we update the weights simultaneously by applying the momentum formula in section 2.2. Please see appendix and [1] for more details on our implementation.

Stochastic gradient descent uses training rate α and momentum m parameters. We initialize $\alpha = 1$ and $m = 0.5$. We decrease α as training progresses and we approach the best solution; specifically, on iteration i, we set $\alpha = \frac{150}{150+i}$. On the other hand, we increase momentum to $m = 0.9$ after twenty iterations once the initial training has mostly stabilized. Preliminary testing shows that these choices yield good performance.

We evaluate our neural net's performance on the Mixed National Institute of Standards and Technology (MNIST) dataset of handwritten digits [9] in terms of overall training algorithm, network structure, number of filters, size of filters, number of pooling regions, and output layer activation function. The MNIST dataset contains 60,000 training images and 10,000 testing images. Its name comes from the NIST dataset that it is an adaptation of. In the NIST dataset, training and testing images were taken from two different sources, while in MNIST, the sources are merged.

5   CNN Evaluation

In this section, we show and analyze CNN performance under various network parameters tested in our experiment. Each parameter takes the default value in Table 1 when unspecified in the subsections below.

Hidden layer configuration	Test error	Time for one iteration (sec)
No hidden layer	2.69%	30.48
10-neuron hidden layer	1.73%	30.80
16-neuron hidden layer	1.74%	30.66
32-neuron hidden layer	1.81%	30.56
32-neuron followed by 16-neuron hidden layers	1.93%	30.43

Table 2: Test error and per-iteration runtime for CNN training with various network configurations.

5.1 CNN performance under different network configurations

To start, we explore how different configurations of the neural network after the convolution and pooling layers affect the performance of our CNN. For each configuration, we trained the network with the 60,000 training samples, and then tested on the 10,000 test samples from MNIST dataset. In each case we train the network for 10 hours.

Table 2 shows that iteration time remains mostly the same. As we go down the table, the model gets more complex, and the test error first goes down and then goes up, suggesting overfitting for the last model.

5.2 Performance with variable filter size

We next investigate how filter size affects the performance of our neural network. We tried out various filter sizes ranging from $n = 5$ to $n = 15$ (for filter sidelength n), again training for ten hours.

The results are shown in Figures 2 and 3. The horizontal axis of both figures show the filter size. For instance, when the filter size is "7 by 7", it achieves an test error rate at about 3.5% and average iteration time takes about 20 seconds.

From the results, we can see that as filter size increases, the average iteration time generally increases. This is probably because low filter size model has fewer number of features, and the update of the features is faster. On the other hand, high filter size model has much more number of features. For instance, the "15 by 15" filter size model has 25 times many features than the "5 by 5" filter size model, and thus the update is much slower, which contributes to slower running time per iteration. From our experiment, the filter size that achieves the fastest running time is "5 by 5",


Figure 2: Test error against filter size.

Figure 3: Per-iteration runtime against filter size.

the smallest of all filter sizes.

The test error rate generally decreases and then increases. This is probably because a very low filter size has too few features, and thus cannot capture all the information from the image; while a very high filter size tends to have too many features, and thus tends to overfit the training data and, as a consequence, performs badly in test data. Another problem with high filter size model is that the training time takes longer due to more features, and thus when terminated at a fixed time (as in our experiment), it will have run much fewer iterations then the low filter size model, i.e. the number of iteration completed is not sufficient. From our experiment, the optimal filter size range is from "9 by 9" to "13 by 13".

Figure 4: Test error against number of filters.

Figure 5: Per-iteration runtime against number of filters.

5.3 Performance with variable filter count

We then investigate how changing the number of filters affects CNN performance. Note that for all experiments in this subsection, we use the default setting of filter size "13 by 13". The number of filters we explored ranges from 8 to 256. Again, we train the network for each filter size for 10 hours.

The results are shown in Figures 4 and 5. The horizontal axis of both figures show the the number of filters used, and is plotted in log scales. For instance, when the number of filters is 16, it achieves an test error rate at about 3% and average iteration time takes about 10 seconds.

From the results, we can see that as the number of filters increases, the average iteration time generally increases, with almost the same speed. This is as expected, because the fewer number of filters in the model, the less complex the model is, and hence the fewer parameters the model needs to update. For instance, the model that with 256 filters have approximately 16 times more parameters (weights, features, etc.) to update compared to the model with 16 filters, and this shows in Figure 5 as the running time per iteration for the former model is about 16 times of that of the latter model.

The test error rate generally decreases and then increases. This is probably due to the same reason that we mention previously: models with low number of filters underfit the training data, while models with high number of filters tends to overfit. From our experiment, the optimal number of filters to use ranges approximately from 16 to 64.

5.4 Performance with variable pooling size

Next we investigate how different pooling sizes affect performance. The number of filters we explored ranges from the extreme case "1 by 1" (a single pooling block that averages over the entire output matrix from the convolution operation) to "6 by 6".

The results are shown in Figures 6 and 7. The horizontal axis of both figures show the the dimension of the pooling size used. For instance, when the pooling size is 4 by 4, it achieves a test error rate at about 4.2% and average iteration time takes about 34 seconds.

From the results, we can see that as the number of pooling size increases, the average iteration time generally increases, with an almost linear relationship. This is as expected, because the fewer number of pooling blocks in the model, the less complex the model is, and hence the fewer parameters the model needs to update. For instance, the model that with pooling sizes of "6 by 6" have many more parameters (specifically weight and bias terms between the pooling layer and the output layer, which are fully connected) to update compared to the model with pooling size "1 by 1", and this shows up in Figure 7.

Again, the test error rate shows a "U-shape" trend, and the reason is similar to that in the previous subsection: models with low number of pooling blocks underfit the training data, while models with high number of pooling blocks tend to overfit. From our experiment, the optimal pooling size is

Figure 6: Test error against pooling size.

Figure 7: Per-iteration runtime against pooling size.

"2 by 2".

5.5 Other parameters

We also experiment with using sigmoid activation function in the output layer (instead of the default softmax activation function used over the other experiments), and training with conventional gradient descent rather than stochastic gradient descent (SGD).

Sigmoid activation resulted in an error rate of 59.62% compared to 2.69% with the default settings for softmax activation function. This indicates that the cross-entropy cost function is critical to training convolutional neural networks for


Figure 8: Training and testing error at different training durations (plotted at fifteen minute intervals).

multi-label classification, otherwise multiple neurons might fire and the squared error cost does not take into account differentiating between the size of the activation outputs.

Gradient descent only completed five iterations in the ten hours allocated for training, and got an error rate of 90.2%. This is equivalent to a random classifier, so clearly the network was not trained for long enough to have a useful structure. Stochastic gradient descent, which achieved the 2.69% error, solves this problem by randomly taking subsets of the data during each iteration so that we can train for a higher number of iterations, even if there is greater variance in the derivatives at each iteration.

Figure 8 shows percent classification error on the training and testing datasets as the CNN training algorithm progresses. At the beginning, the improvement is much more significant, while later on we approach the best solution and are no longer able to make significant gains; in the later three hours shown, we only decrease the error rate by an additional few percentage points. Also, as a result of the stochastic gradient algorithm (which updates based on random samples at each iteration), the training error is not monotonically decreasing (although it is still surprising that there are such large variations in the errors even while the network is still training overall; each fifteen-minute interval in the plot includes thirty or so SGD iterations).

Figure 9 shows four examples convolutional filters after training. It appears that these filters recognize parts of digits. Black portions indicate positive values, which match up with marks in the image, while white portions indicate negative values (the sign does not actually matter since it may be flipped by a later weight). Thus, each filter

Figure 9: Examples of convolutional filters. These filters are 13 by 13 pixels. The top two filters are from a 32-filter CNN while the bottom two are from an 8-filter CNN. White indicates negative filter values while black indicates positive values.

not only recognizes certain edges or combinations of edges in the image, but also avoids cases where other edges are present; for example, in the lower left filter, there is a black horizontal line through the center but also a white vertical line along the left side. Once filters have been convolved with the image and pooled across blocks, the output layer weights can take care of piecing the parts together to recognize the entire handwritten digit.

5.6 Heatmap

In the experiments above, we have evaluated the effect of varying one parameter on performance. Here, we take a step further by trying experimentations that vary *two* parameters simultaneously. This means for the pooling size, filter size, and filter count parameters, we keep one parameter fixed at default value, and test different pairs of the other two parameters. Due to time constraints on training, we only train each CNN for three hours. The result is presented in heat map below, with different colors showing different error rates.

Figures 10 and 11 show how varying pooling size and filter size affects the performance of our neural net model. Figure 10 shows a heatmap of training error against pooling size and filter size, while Figure 11 shows a heatmap of testing error against


Figure 10: Heatmap of training error against pooling size and filter size.

Figure 11: Heatmap of testing error against pooling size and filter size.

pooling size and filter size. From the results, we can see that the training error is generally smaller than the testing error, as what we have expected. Moreover, we see that there is a strong correlation between training error and testing error, i.e. model with parameters that does a bad job in the training set also performs badly in the testing set. We additionally observe that the optimal neural net model has pooling size ranging from about "1 by 1" to "3 by 3", and filter size ranging from "11 by 11" to "15 by 15". This is probably because with the parameters above, the CNN training not only runs faster, meaning it is able to perform sufficient iterations in three hours, but also has enough complexity to describe the underlying data.

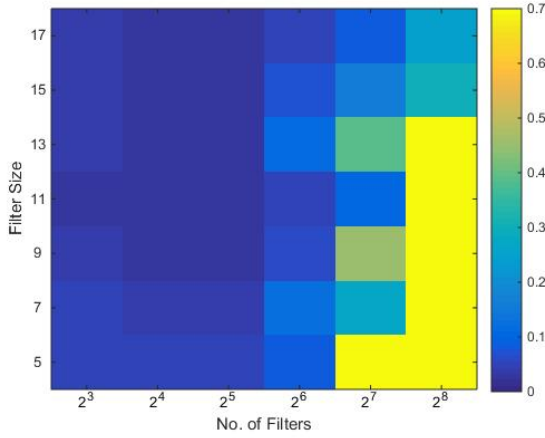Figures 12 and 13 show how varying filter count

Figure 12: Heatmap of training error against filter size and filter count.

Figure 13: Heatmap of training error against filter size and filter count.

and filter size affects the performance of our neural net model. Figure 12 shows a heatmap of training error against filter count and filter size, while Figure 13 shows a heatmap of testing error against filter count size and filter size. From the heat maps, we again conclude that the training error is generally smaller than the testing error, and that there is a strong correlation between training error and testing error. We observe that the optimal neural net model has filter count ranging from about 8 to 32, and filter size ranging from "9 by 9" to "17 by 17". Similar to with pooling size and filter size, this is probably because with these parameters, the neural net model both trains faster and has sufficient complexity.

6   Comparison with HOG-feature recognition

Before deep learning techniques became widely used in image recognition, various other methods were used to extract the features from the image and models were trained on these features. Examples include histogram of oriented gradients (HOG) and scale-invariant feature transform (SIFT). In this section, we will compare the performance of our CNN implementation with the performance of using HOG with SVM classifier, using different percentages of the 60,000 training samples. We use the default parameters except we include a ten-neuron fully connected hidden layer between the pooling and output layers.

We implemented the CNN library for this project, but we used MATLAB packages for the HOG feature extraction method. HOG divides an image into cells and computes a local histogram of gradient orientations in each cell. The histograms and normalized and grouped across block regions, and these descriptor values can then be fed into a classifier.

The results are shown in Figures 14 and 15. From the figures, we can see that the training error rate is generally lower than the testing error rate, and this is expected because our model is based on the training set. Moreover, as we use more training data, we generally get a higher training error, but a lower testing error. This is because with less training data, it is much easier to fit a model that does very well on the training set, but the same model does not generalize well to the overall data, and thus performs badly on the testing set. As the training set data increases, although we get a higher training error as it becomes increasingly harder to fit a perfect model on the training set, the testing error generally decreases. This agrees with our intuition.

Moreover, we can see that our CNN performs much better than the HOG feature-extraction technique. At any percentage of training data, our model achieves a test error strictly less than the HOG feature-extraction technique. When all training data is used, our neural net model achieves an error rate of about 2.25 percent test error, while the HOG feature-extraction technique achieves more than 5 percent. This perhaps explains why the old feature-extraction techniques, such as HOG and SIFT, are becoming increasingly obsolete today, due to the advent of deep learning.
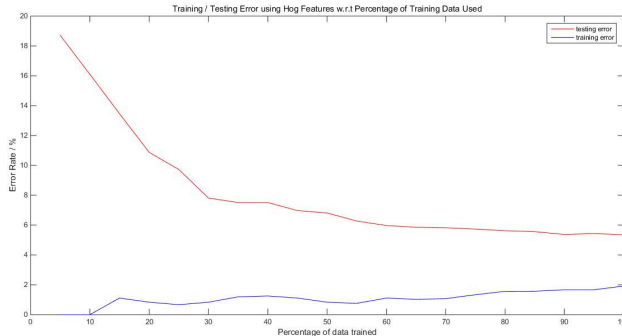
Figure 14: Training and testing error using HOG features against percentage of training data used (misclassification percentage rate on vertical axis, percentage of training data on horizontal axis; testing error in red and training error in blue).


Figure 15: Training and testing error of CNN against percentage of training data used (misclassification percentage rate on vertical axis, percentage of training data on horizontal axis; testing error in red and training error in blue).

7   Conclusion

With our basic CPU-based Python convolutional neural network toolkit, we are able to achieve error rates as low as 1.73% for certain network parameter choices. Deep learning algorithms like CNNs allow machine learning of the very features to be used for classification; this contrasts with previous approaches where feature functions are designed by humans. CNNs show that biologically inspired machine learning algorithms often can improve performance; this echoes previous work in genetic algorithms and ant colony optimization.

However, CNN training did take significantly longer than feature extraction and SVM classification. Still, GPU-based CNN libraries take advantage of parallelizable sequences of matrix operations that perform very well on modern GPUs. Additionally, optimizations can be made when using stochastic gradient descent for the batch size.

There were also cases where we used a CNN that was too complex (too many filters or too many network layers) and resulted in overfitting. There has been significant research in preventing CNN overfitting. One approach called "dropout" randomly disables half of the filters on each training iteration and yields promising performance [3].

Another area where we did not get to investigate in greater detail is using the features from the convolutional and pooling layers to train an SVM classifier. Once we have trained the filters (and corresponding bias terms), we can remove the following layers of the CNN and simply feed the output of the pooling layer for all training and test samples into the classifier; this provides an alternative for the classification component, even though we still use convolution and pooling operations for learning features. So, a future comparison between SVM with these features and the full convolutional neural network would be interesting.

8   Appendix: sample source code

We include samples of source code from our CNN library and plotting routines below. The full code is available in a git repository that can be accessed at [1].

8.1   CNN library (Python)

```python
class BasicNetwork:
  def __init__(self, shape):
```

```python
        # shape is an iterable of integers, each
            representing a layer
        # each integer specifies the number of
            neurons in that layer
        self.nlayers = len(shape)
        self.layers = []
        self.shape = list(shape)

        # first layer is input layer
        self.layers.append({'type': 'input'})

        # process other layers
        for i in xrange(len(shape) - 1):
            if isinstance(shape[i + 1], dict):
                if shape[i + 1]['type'] ==
                    'convsample':
                    # add a 2D
                        convolutional/subsampling
                        layer
                    m = shape[i + 1]['m'] # input
                        sidelength
                    c = shape[i + 1]['c'] # number
                        of channels in image
                    n = shape[i + 1]['n'] # filter
                        sidelength
                    k = shape[i + 1]['k'] # number
                        of convolutional filters
                    p = shape[i + 1]['p'] # pooled
                        region length
                    conv_count = m - n + 1 #
                        convolution output
                        sidelength

        ....

    def forward(self, inputs,
        return_activations = False):
        # inputs specifies one float for each
            neuron in the first layer
        # if return_activations is true, we
            return a list of activations at
            each layer
        # otherwise, we only return the output
            layer
        values = numpy.array([inputs], float).T
            # values in current layer
        activations = [{'activations': values}]

        for layer in self.layers[1:]:
            if layer['type'] == 'sigmoid':
                # compute the weighted sum of
                    neuron inputs, plus bias term
                    (for each neuron in current
                    layer)
                z_vector =
                    numpy.dot(layer['weights'],
                    values) + layer['bias']

                # apply sigmoid activation function
                values = util.sigmoid(z_vector)
                if return_activations:
                    activations.append({'activations':
                        values[:, 0]})
            elif layer['type'] == 'softmax':
                .....


    def backward(self, inputs,
        desired_outputs):
        # inputs is a list of input layer
            values, desired_outputs is expected
            output layer values
        inputs = numpy.array([inputs], float).T
        desired_outputs =
            numpy.array([desired_outputs],
            float).T

        # execute forward propogation and store
            activations of each layer
        activations = self.forward(inputs[:, 0],
            True)

        # compute deltas at each layer
        deltas_list = [None] * self.nlayers #
            deltas_list[0] is for input layer
            and remains None

        if self.layers[self.nlayers - 1]['type']
            == 'sigmoid':
            cost =
                numpy.sum(abs(activations[self.nlayers
                - 1]['activations'] -
                desired_outputs[:, 0]))
            deltas_list[self.nlayers - 1] =
                numpy.multiply(activations[self.nlayers
                - 1]['activations'] -
                desired_outputs[:, 0],
                util.sigmoid_d2(activations[self.nlayers
                - 1]['activations']))
        elif self.layers[self.nlayers -
            1]['type'] == 'softmax':
            cost =
                -numpy.sum(numpy.multiply(desired_outputs[:,
                0],
                numpy.log(activations[self.nlayers
                - 1]['activations'])))
            deltas_list[self.nlayers - 1] =
                activations[self.nlayers -
                1]['activations'] -
                desired_outputs[:, 0]
        else:
            raise Exception('invalid error
                function type')

        for l in xrange(self.nlayers - 2, 0,
            -1): # for each non-input layer
            previous_deltas = deltas_list[l + 1]
            target_layer = self.layers[l + 1]
```

```python
        if target_layer['type'] == 'sigmoid'
          or target_layer['type'] ==
          'softmax':
            sums =
                numpy.dot(target_layer['weights'].T,
            .....
```

8.2 Plotting routines (MATLAB)

```matlab
no_of_filter_vec = 2.^[3 4 5 6 7 8];
testing_error_vec = [334 304 266 269 364
    626]/10000;
%plot(no_of_filter_vec,testing_error_vec,'-b')
plot (log2(no_of_filter_vec),
    testing_error_vec)

set(gca,'XTick',[2:9])
set(gca, 'XTickLabel',[])
xt = get(gca, 'XTick');
yl = get(gca, 'YLim');
str = cellstr( num2str(xt(:),'2^{%d}') ); %#
    format x-ticks as 2^{xx}
hTxt = text(xt, yl(ones(size(xt))), str, ...
    %# create text at same locations
    'Interpreter','tex', ...              %#
        specify tex interpreter
    'VerticalAlignment','top', ...        %#
        v-align to be underneath
    'HorizontalAlignment','center');      %#
        h-aligh to be centered
ylabel('test error')
xlabel(sprintf('\nno. of filters'))


no_of_filter_vec = 2.^[3 4 5 6 7 8];
avg_iter_time = 10*60*60./[9659 4563 2343
    1181 592 312];
%plot(no_of_filter_vec,testing_error_vec,'-b')
plot (log2(no_of_filter_vec), avg_iter_time)

set(gca,'XTick',[2:9])
set(gca, 'XTickLabel',[])
xt = get(gca, 'XTick');
yl = get(gca, 'YLim');
str = cellstr( num2str(xt(:),'2^{%d}') ); %#
    format x-ticks as 2^{xx}
hTxt = text(xt, yl(ones(size(xt))), str, ...
    %# create text at same locations
    'Interpreter','tex', ...              %#
        specify tex interpreter
    'VerticalAlignment','top', ...        %#
        v-align to be underneath
    'HorizontalAlignment','center');      %#
        h-aligh to be centered
ylabel('test error')
```

```matlab
xlabel(sprintf('\nno. of filters'))

filter_size_vec = [5 7 9 11 13 15];
testing_error_vec = [605 369 301 359 269
    406]/10000;
plot(filter_size_vec,testing_error_vec,'-b')
xlabel('filter size')
ylabel('test error')

filter_size_vec = [5 7 9 11 13 15];
avg_iter_time = 10*60*60 ./ [2290 1744 1521
    1350 1181 1227];
plot(filter_size_vec,avg_iter_time,'-b')
xlabel('filter size')
ylabel('time of one iteration / sec')
```

References

[1] Favyen Bastani and Zhengli Wang. CNN git repository, December 2014. https://github.com/uakfdotb/cnn.

[2] Andrea Frome, German Cheung, Ahmad Abdulkader, Marco Zennaro, Bo Wu, Alessandro Bissacco, Hartwig Adam, Hartmut Neven, and Luc Vincent. Large-scale privacy protection in google street view. In *IEEE International Conference on Computer Vision*, 2009.

[3] Geoffrey E Hinton, Nitish Srivastava, Alex Krizhevsky, Ilya Sutskever, and Ruslan R Salakhutdinov. Improving neural networks by preventing co-adaptation of feature detectors. *CoRR*, abs/1207.0580, 2012.

[4] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E. Hinton. Imagenet classification with deep convolutional neural networks. In F. Pereira, C.J.C. Burges, L. Bottou, and K.Q. Weinberger, editors, *Advances in Neural Information Processing Systems 25*, pages 1097–1105. Curran Associates, Inc., 2012.

[5] LISA Lab. DeepLearning documentation: Convolutional nerual networks (LeNet), December 2014. http://deeplearning.net/tutorial/lenet.html.

[6] Stanford Deep Learning Lab. Optimization: Stochastic gradient descent. http://ufldl.stanford.edu/tutorial/supervised/OptimizationStochasticGradientDescent/.

[7] Steve Lawrence, C Lee Giles, Ah Chung Tsoi, and Andrew D Back. Face recognition: A convolutional neural-network approach. *IEEE*

Transactions on Neural Networks, 8(1):98–113, 1997.

[8] Patrick Le Callet, Christian Viard-Gaudin, and Dominique Barba. A convolutional neural network approach for objective video quality assessment. *IEEE Transactions on Neural Networks*, 17(5):1316–1327, 2006.

[9] Yann LeCun, Corinna Cortes, and Christopher J.C. Burges. MNIST database of handwritten digits. http://yann.lecun.com/exdb/mnist/.

[10] Genevieve B. Orr. Neural network classification. http://www.willamette.edu/~gorr/classes/cs449/.